



GPT-2 Heads — Figure 1 analogue; sample 0 from N=1000


L03H07: periods/newlines
shown-key mass=0.73




L11H04: boundary candidate
shown-key mass=0.84



L08H03: initial-key sink
shown-key mass=0.98



L05H12: previous-token attention
shown-key mass=0.96



L01H12: broad attention
shown-key mass=0.40

